

Los Angeles R Users' Group

---

# NumPy and SciPy for Data Mining and Data Analysis Including iPython, SciKits, and matplotlib

Ryan R. Rosario

---

March 29, 2011

# What is NumPy?



NumPy is the standard package for scientific and numerical computing in Python:

- powerful  $N$ -dimensional array object.
- sophisticated (broadcasting) functions.
- tools for integrating C/C++ and Fortran code (R is similar)
- useful linear algebra, Fourier transform and random number capabilities.

Additionally...

- data structures can be used as efficient multi-dimensional container of arbitrary data types.

# So, What is SciPy?



Often (incorrectly) used as a synonym of NumPy.

- SciPy *depends on* NumPy's data structures.
- OSS for mathematics, science, and engineering.
- provides efficient numerical routines for numerical integration and optimization.

“NumPy provides the data structures, SciPy provides the application layer.”

# Getting NumPy

The ease of installation is highly dependent on OS and CPU.

- **Official** releases are recommended and available for
  - ① 32 bit MacOS X
    - **Caveat:** Must use Python Python, not Apple Python.
  - ② 32 bit Windows
- **Unofficial**/third-party binary releases for
  - ① All Linux (use platform package system)
  - ② 64 bit MacOS X (SciPy Superpack)
    - Requires Apple Python 2.6.1
  - ③ 64-bit Windows (Christoph Gohlke, UC Irvine)
- Source Code
- Enthought is to SciPy/Numpy as Revolution is to R.

Main download link is <http://www.scipy.org/Download>

# Disclaimer

Time is even more limited than usual, so the goal is to illustrate what NumPy is, and what NumPy, SciPy and its colleagues can achieve. Some goals:

- Brief NumPy tutorial via demo.
- Uses of NumPy and SciPy, including SciKits.
- Brief data analysis demo with NumPy, SciPy and matplotlib.
- Concluding remarks.

# NumPy Basics

The main data structure is a *homogeneous multidimensional array* – a table of elements **all of the same type**, indexed by a tuple of positive integers. Some examples are vectors, matrices, images and spreadsheets.

**Why not just use a list of lists?** NumPy provides functions that operate on an entire array at once (like R), Python does not...really.

NumPy is full of syntax and time is limited, so let's jump straight into a demo...

# NumPy Basics

Demo

## Keep in Mind: Broadcasting vs. Recycling

In NumPy, if two arrays do not have the same length and we try to, say, add them, 1 is appended to the smaller array so that the dimensions match.

In R, if two vectors do not have the same length, the values in the shorter vector are recycled and a warning is generated. Personally, recycling seems more practical.



# Keep in Mind: Indexing!

R starts its indices at 1 and ranges are inclusive (i.e. `seq(1,3)` -> 1, 2, 3)

NumPy (and Python) starts its indices at 0 and ranges are exclusive (i.e. `arange(1,3)` -> 1, 2).

# Keep in Mind: Data Type Matters!

R does a good job of using the correct data type.

NumPy requires you to make decisions about data type. Using `int64` for 1 is typically wasteful.

ipython: a replacement CLI I

IP[y]:

IPython is the recommended CLI for NumPy/SciPy. I will use IPython in my demo.

- 1 an enhanced interactive Python shell.
- 2 an architecture for interactive parallel computing.

# ipython: a replacement CLI II

Some advantages over the standard CLI:

- ① tab completion for object attributes and filenames, auto parentheses and quotes for function calls.
- ② simple object querying
  - `object_name?` prints details about an object.
  - `%pdoc` (docstring), `%pdef` (function definition), `%psource` (full source code).
- ③ `%run` to run a code file, much like `source` in R. Differs from `import` because the code file is reread each time `%run` is called. Also provides special flags for profiling and debugging.

## ipython: a replacement CLI III

- ④ `%pdb` for spawning a `pdb` debug session.
- ⑤ **Output cache** saves the results of executing each line of code. `_`, `__`, `---` save the last three results.
- ⑥ **Input cache** allows re-execution of previously executed code using a snapshot of the original input data in variable `In`. (to reexecute lines 22 through 28, and line 34: `exec In[22:29] + In[34]`).
- ⑦ **Command history** with(out) line numbers using `%hist`. Can also toggle journaling with `%logstart`.
- ⑧ Execute system commands from within IPython as if you are in the terminal by prefixing with `!`. Use `!!`, `%sc`, `%sx` to get output into a Python variable (`var = !ls -la`). Use Python variables when calling the shell.

## ipython: a replacement CLI IV

- ⑨ Easy access to the profiler, `%prun`.
- ⑩ Easy to integrate with `matplotlib`.
- ⑪ Custom profiles (options, modules to load, function definitions). Similar to R workspace.
- ⑫ Run other programs (that use pipes) and extract output (`gnuplot`).
- ⑬ Easily run doctests.
- ⑭ Lightweight version control.
- ⑮ Save/restore session, like in R (the authors specifically said “like R”).


# matplotlib: SciPy's Graphics System I



2D graphics library for Python that produces publication quality figures. Quality with respect to R is in the eye of the beholder.


- Graphics in R are much easier to construct, and require less code.
- Graphics in `matplotlib` may look nicer; require much more code.

# matplotlib: SciPy's Graphics System II





# matplotlib: SciPy's Graphics System III



# matplotlib: SciPy's Graphics System IV

Show matplotlib gallery.

I will do a demo of matplotlib in the data analysis demo, if time.

# Quick Linear Regression Demo

Demo here.

# SciKits I

SciKits are modules or “plug-ins” for SciPy; they are too specialized to be included in SciPy itself. Below is a list of some relevant SciKits.

- ❶ statsmodels: OLS, GLS, WLS, GLM, M-estimators for robust regression.
- ❷ timeseries: time series analysis.
- ❸ bootstrap: bootstrap error-estimation.
- ❹ learn: machine learning and data mining

# SciKits II

- ⑤ `sparse`: sparse matrices
- ⑥ `cuda`: Python interface to GPU powered libraries.
- ⑦ `image`: image processing
- ⑧ `optimization`: numerical optimization

For more information, see


<http://scikits.appspot.com/scikits>.

# Scientific Software Relying on SciPy I

Many exciting software packages rely on SciPy



**Python Based Reinforcement Learning, Artificial Intelligence and Neural Network Library.** Built around a neural networks kernel. Provides unsupervised learning and evolution.  
<http://www.pybrain.org>



**PyMC** Pythonic Markov Chain Monte Carlo. MCMC, Gibbs sampling, diagnostics



**PyML:** SVM, nearest neighbor, ridge regression, multiclass methods, model selection, filtering, CV, ROC.

# Scientific Software Relying on SciPy II



**Monte:** gradient based learning machines, logistic regression, conditional random fields.



**MILK:** SVM, k-NN, random forests, decision trees, feature selection, affinity propagation.



**Modular toolkit for Data Processing:** signal processing, PCA, ICA, manifold learning, factor analysis, regression, neural networks, mixture models.

# Scientific Software Relying on SciPy III



**PyEvolve:** Genetic algorithms

theano

**Theano:** Theano is a Python library that allows you to define, optimize, and evaluate mathematical expressions involving multi-dimensional arrays efficiently: GPUs, symbolic expressions, dynamic C code generation.



**Orange:** Data mining, visualization and machine learning. Full GUI and Python scripting.



# Scientific Software Relying on SciPy IV



**SAGE:** a hodgepodge of open-source software designed to be one analysis platform.

# Pros/Cons I

What NumPy/SciPy does better:

- ➊ Much larger community (SciPy + Python).
- ➋ Better, and cleaner graphics...slightly.
- ➌ **Sits on atop a programming language rather than being one.**
- ➍ Dictionary data type is very useful (Python).
- ➎ Adheres to a system of thought (“Pythonic”) that is more flexible.
- ➏ Memory mapped files and extensive sparse matrix packages.
- ➐ Easier to incorporate into a company/development workflow.
- ➑ **Opinion:** Developers don’t have to dive into C as soon as they do with R.

# Pros/Cons II

What R does better:

- ➊ One word `data.frame`.
- ➋ Installation is simple and universal.
- ➌ Application is self-contained (NumPy consists of several moving parts).
- ➍ Very simple, user friendly syntax.
- ➎ Graphics and analysis are fairly quick to code.
- ➏ Much easier data input.
- ➐ Handles missing values better.

The SciPy community is very useful, prolific and friendly. Project has very high morale. Some places to get help:

- ➊ <http://www.numpy.org>: extensive documentation, examples and cookbook.
- ➋ <http://www.scipy.org>: extensive documentation, examples and cookbook.
- ➌ NumPy and SciPy mailing lists are helpful.
- ➍ StackOverflow
- ➎ IRC channel #scipy
- ➏ Frequent sprints and hackathons.
- ➐ SciPyCon (Austin, TX)
- ➑ SoCalPiggies and LA Python (first meeting April 1!)

# Keep in Touch!

My email: `ryan@stat.ucla.edu`

My blog: `http://www.bytemining.com`

Follow me on Twitter: `@datajunkie`

# References

- ① NumPy <http://www.numpy.org>
- ② SciPy <http://www.scipy.org>
- ③ SciKits <http://scikits.appspot.com>
- ④ iPython <http://ipython.scipy.org/moin/>
- ⑤ matplotlib <http://matplotlib.sourceforge.net/>